

# Big Data Analytics

## # 9: Efficient In-Memory Analytics with Polars

Instructor: Oleh Tymchuk

# #11: Agenda

- Introduction to Polars
- Polars vs. Other Tools
- Use cases
- Practical cases
- Useful Links

# Introduction to Polars

# What is Polars?

## Definition:

- High-performance DataFrame library for Python and Rust
- Built on Apache Arrow for memory efficiency
- Written in Rust, supports parallel execution

## Key Comparisons:

- Faster than Pandas for large datasets
- Similar speed to Spark/Dask, but simpler—no cluster setup

# Why Use Polars?

## **Speed**

5x–50x faster than Pandas in many workloads

## **Memory Efficiency**

Uses Apache Arrow's efficient columnar memory format

## **Lazy Execution**

Defers computation for optimization

## **Clean API**


Chainable methods for readable, maintainable code

# Key Features of Polars

Feature	Description
Lazy Evaluation	Optimize complex queries automatically
Arrow Integration	Zero-copy interoperability with Arrow
Parallel Execution	Utilizes all available CPU cores
Rich Data Types	Supports nested data, nulls
Interoperability	Reads/Writes CSV, Parquet, Pandas, NumPy

## Polars vs. Other Tools


# Duration in Seconds






# Data read from disk (Parquet)

Runtime including data read from disk (Parquet)  
(lower is better)



# Data read from disk (CSV)

Runtime including data read from disk (CSV)  
(lower is better)



# Use Cases

# Use Cases

- Large datasets: Process GBs of data on a single machine.
- ETL pipelines: Fast transformations and aggregations.
- Data analysis: Interactive workflows with lazy evaluation.
- Arrow ecosystems: Integrate with tools like DuckDB, PySpark.

# Limitations

- Smaller community vs. Pandas
- Fewer third-party integrations
- Learning curve for expression-based syntax

# Practical cases

# Useful Links

[Python API reference](#)

[Polars Academy](#)

Q&A